

# hw1

February 2, 2018

## 1 Homework 1

*This notebook includes both coding and written questions. Please hand in this notebook file with all the outputs and your answers to the written questions.*

This assignment covers linear filters, convolution and correlation

```
In [2]: # Setup
import numpy as np
import matplotlib.pyplot as plt
from time import time
from skimage import io

from __future__ import print_function

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
%load_ext autoreload
%autoreload 2

print("Huynh Bao Quoc (B1400516)")
print("Nguyen Trung Hau (B1400493)")
```

Huynh Bao Quoc (B1400516)

Nguyen Trung Hau (B1400493)

### 1.1 Part 1: Convolutions

#### 1.1.1 1.1 Commutative Property (10 points)

Recall that the convolution of an image  $f : \mathbb{R}^2 \rightarrow \mathbb{R}$  and a kernel  $h : \mathbb{R}^2 \rightarrow \mathbb{R}$  is defined as follows:

$$(f * h)[m, n] = \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} f[i, j] \cdot h[m - i, n - j]$$

Or equivalently,

$$(f * h)[m, n] = \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} h[i, j] \cdot f[m - i, n - j] \quad (1)$$

$$= (h * f)[m, n] \quad (2)$$

Show that this is true (i.e. prove that the convolution operator is commutative:  $f * h = h * f$ ).

Chng minh rng phép tích chp có tính giao hoán:  $f * h = h * f$ .

**Your Answer:** Write your solution in this markdown cell. Please write your equations in [LaTeX equations](#).

$$V1 : (f * h)[m, n] = \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} f[i, j] \cdot h[m - i, n - j]$$

$t :$

$$u = m - i$$

$$v = n - j$$

Suyra :

$$i \rightarrow -\infty \Rightarrow u \rightarrow \infty, j \rightarrow -\infty \Rightarrow v \rightarrow \infty$$

$$i \rightarrow \infty \Rightarrow u \rightarrow -\infty, j \rightarrow \infty \Rightarrow v \rightarrow -\infty$$

$$\Leftrightarrow (f * h)[m, n] = \sum_{u=-\infty}^{\infty} \sum_{v=-\infty}^{\infty} f[m - u, n - v] \cdot h[u, v]$$

$$\text{Thay : } u = i, v = j$$

$$V2 : (f * h)[m, n] = \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} f[m - i, n - j] \cdot h[i, j]$$

### 1.1.2 1.2 Linear and Shift Invariance (10 points)

Let  $f$  be a function  $\mathbb{R}^2 \rightarrow \mathbb{R}$ . Consider a system  $f \xrightarrow{S} g$ , where  $g = (f * h)$  with some kernel  $h : \mathbb{R}^2 \rightarrow \mathbb{R}$ . Show that  $S$  defined by any kernel  $h$  is a Linear Shift Invariant (LSI) system. In other words, for any  $h$ , show that  $S$  satisfies both of the following:

- If  $f[m, n] \xrightarrow{S} g[m, n]$  then  $f[m - m_0, n - n_0] \xrightarrow{S} g[m - m_0, n - n_0]$

Cho  $f$  là mt hàm  $\mathbb{R}^2 \rightarrow \mathbb{R}$ . Xét mt h thng x lý  $f \xrightarrow{S} g$ , trong ó  $g = (f * h)$  vi mt kernel (mt n chp)  $h : \mathbb{R}^2 \rightarrow \mathbb{R}$ . Hãy chng minh rng h thng  $S$  vi kernel  $h$  bt k là mt h thng Linear Shift Invariant (LSI). Nói cách khác, vi  $h$  bt k, hãy ch ra rng  $S$  tho c hai iu kin sau:

- If  $f[m, n] \xrightarrow{S} g[m, n]$  then  $f[m - m_0, n - n_0] \xrightarrow{S} g[m - m_0, n - n_0]$

**Your Answer:** Write your solution in this markdown cell. Please write your equations in [LaTeX equations](#).

Nhp câu tr li ca bn ây.

### 1.1.3 1.3 Implementation (30 points)

In this section, you will implement two versions of convolution: - conv\_nested - conv\_fast  
First, run the code cell below to load the image to work with.

```
In [18]: # Open image as grayscale
img = io.imread('dog.jpg', as_grey=True)

# Show image
plt.imshow(img)
plt.axis('off')
plt.title("Isn't he cute?")
plt.show()
```

Isn't he cute?



Now, implement the function `conv_nested` in `filters.py`. This is a naive implementation of convolution which uses 4 nested for-loops. It takes an image  $f$  and a kernel  $h$  as inputs and outputs the convolved image ( $f * h$ ) that has the same shape as the input image. This implementation should take a few seconds to run.

- Hint: It may be easier to implement  $(h * f)$

We'll first test your `conv_nested` function on a simple input.

```
In [19]: from filters import conv_nested
```

```
# Simple convolution kernel.
kernel = np.array(
    [
        [1,0,1],
        [0,0,0],
        [1,0,1]
    ])

# Create a test image: a white square in the middle
test_img = np.zeros((9, 9))
test_img[3:6, 3:6] = 1

# Run your conv_nested function on the test image
test_output = conv_nested(test_img, kernel)

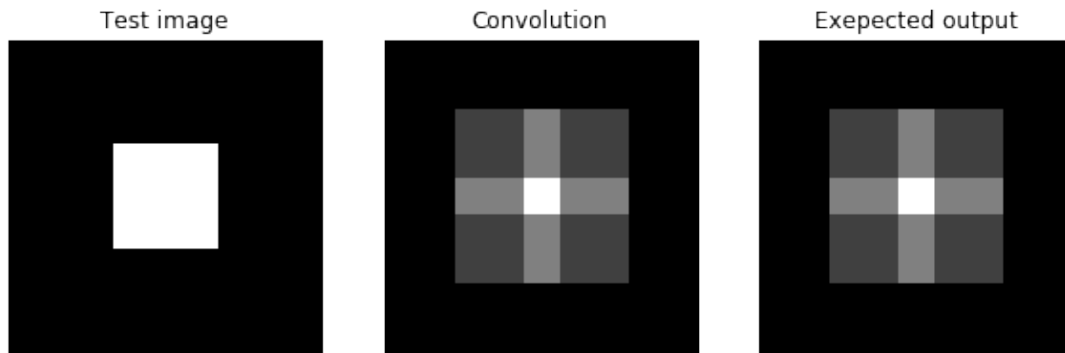
# Build the expected output
expected_output = np.zeros((9, 9))
expected_output[2:7, 2:7] = 1
expected_output[4, 2:7] = 2
expected_output[2:7, 4] = 2
expected_output[4, 4] = 4

# Plot the test image
plt.subplot(1,3,1)
plt.imshow(test_img)
plt.title('Test image')
plt.axis('off')

# Plot your convolved image
plt.subplot(1,3,2)
plt.imshow(test_output)
plt.title('Convolution')
plt.axis('off')

# Plot the expected output
plt.subplot(1,3,3)
plt.imshow(expected_output)
plt.title('Expected output')
plt.axis('off')
plt.show()

# Test if the output matches expected output
assert np.max(test_output - expected_output) < 1e-10, "Your solution is not correct."
```



Now let's test your `conv_nested` function on a real image.

```
In [20]: from filters import conv_nested

# Simple convolution kernel.
# Feel free to change the kernel and to see different outputs.
kernel = np.array(
[
    [1,0,-1],
    [2,0,-2],
    [1,0,-1]
])

out = conv_nested(img, kernel)

# Plot original image
plt.subplot(2,2,1)
plt.imshow(img)
plt.title('Original')
plt.axis('off')

# Plot your convolved image
plt.subplot(2,2,3)
plt.imshow(out)
plt.title('Convolution')
plt.axis('off')

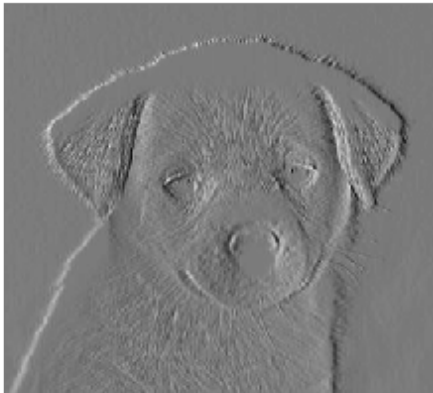
# Plot what you should get
solution_img = io.imread('convoluted_dog.jpg', as_grey=True)
plt.subplot(2,2,4)
plt.imshow(solution_img)
plt.title('What you should get')
plt.axis('off')
```

```
plt.show()
```

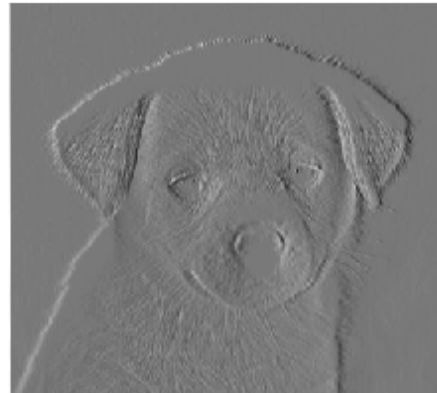
Original



Convolution



What you should get



Let us implement a more efficient version of convolution using array operations in numpy. As shown in the lecture, a convolution can be considered as a sliding window that computes sum of the pixel values weighted by the flipped kernel. The faster version will i) zero-pad an image, ii) flip the kernel horizontally and vertically, and iii) compute weighted sum of the neighborhood at each pixel.

First, implement the function `zero_pad` in `filters.py`.

```
In [5]: from filters import zero_pad
```

```
pad_width = 20 # width of the padding on the left and right
pad_height = 40 # height of the padding on the top and bottom

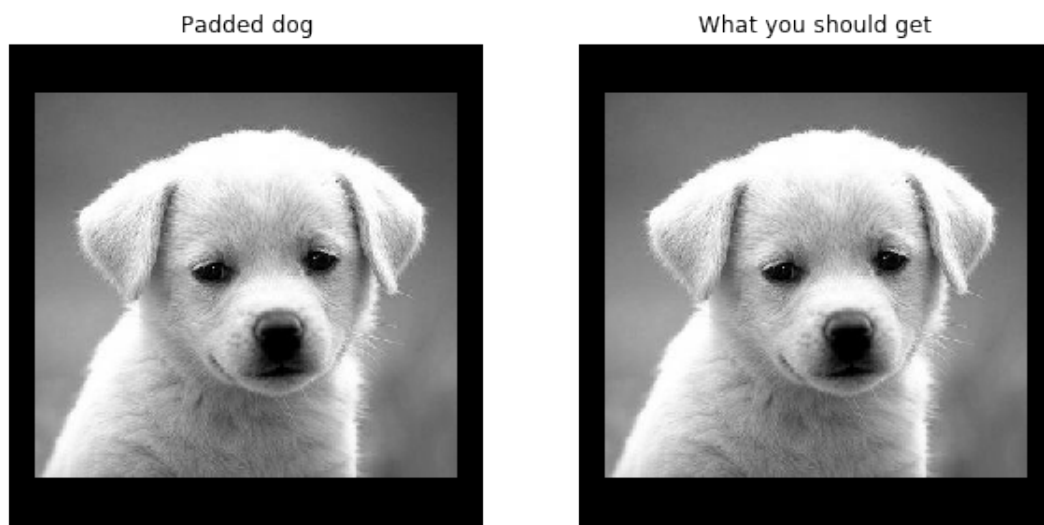
padded_img = zero_pad(img, pad_height, pad_width)

# Plot your padded dog
plt.subplot(1,2,1)
```

```
plt.imshow(padded_img)
plt.title('Padded dog')
plt.axis('off')

# Plot what you should get
solution_img = io.imread('padded_dog.jpg', as_grey=True)
plt.subplot(1,2,2)
plt.imshow(solution_img)
plt.title('What you should get')
plt.axis('off')

plt.show()
```



Next, complete the function `conv_fast` in `filters.py` using `zero_pad`. Run the code below to compare the outputs by the two implementations. `conv_fast` should run significantly faster than `conv_nested`.

Depending on your implementation and computer, `conv_nested` should take a few seconds and `conv_fast` should be around 5 times faster.

```
In [6]: from filters import conv_fast

t0 = time()
out_fast = conv_fast(img, kernel)
t1 = time()
out_nested = conv_nested(img, kernel)
t2 = time()

# Compare the running time of the two implementations
print("conv_nested: took %f seconds." % (t2 - t1))
```

```

print("conv_fast: took %f seconds." % (t1 - t0))

# Plot conv_nested output
plt.subplot(1,2,1)
plt.imshow(out_nested)
plt.title('conv_nested')
plt.axis('off')

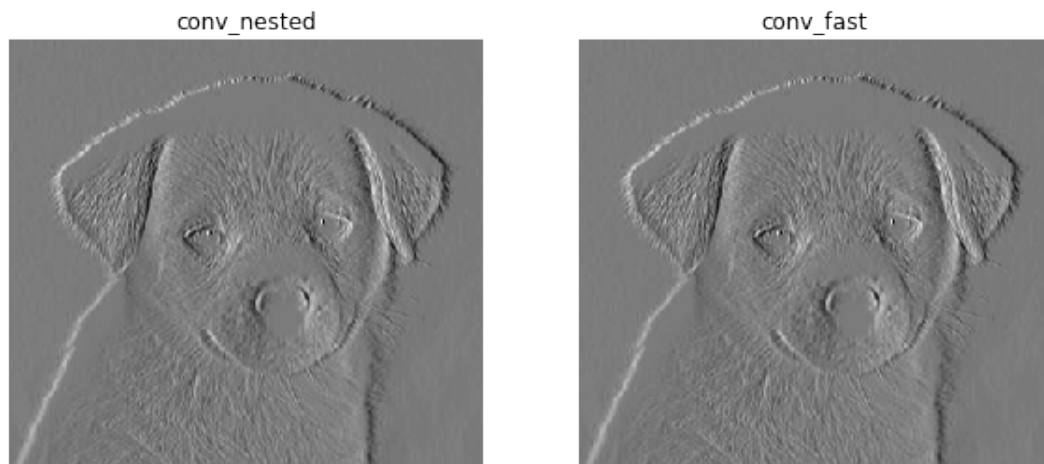
# Plot conv_fast output
plt.subplot(1,2,2)
plt.imshow(out_fast)
plt.title('conv_fast')
plt.axis('off')

# Make sure that the two outputs are the same
if not (np.max(out_fast - out_nested) < 1e-10):
    print("Different outputs! Check your implementation.")

```

conv\_nested: took 1.054778 seconds.

conv\_fast: took 0.505025 seconds.



#### 1.1.4 Extra Credit 1 (1% of final grade)

Devise a faster version of convolution and implement `conv_faster` in `filters.py`. You will earn extra credit only if the `conv_faster` runs faster (by a fair margin) than `conv_fast` **and** outputs the same result.

```

In [7]: from filters import conv_faster
        kernel = np.array(

```



```

[
    [1,0,-1],
    [2,0,-2],
    [1,0,-1]
])

t0 = time()
out_fast = conv_fast(img, kernel)
t1 = time()
out_faster = conv_faster(img, kernel)
t2 = time()
out_fast=out_fast-(out_fast-out_faster)

# Compare the running time of the two implementations
print("conv_fast: took %f seconds." % (t1 - t0))
print("conv_faster: took %f seconds." % (t2 - t1))

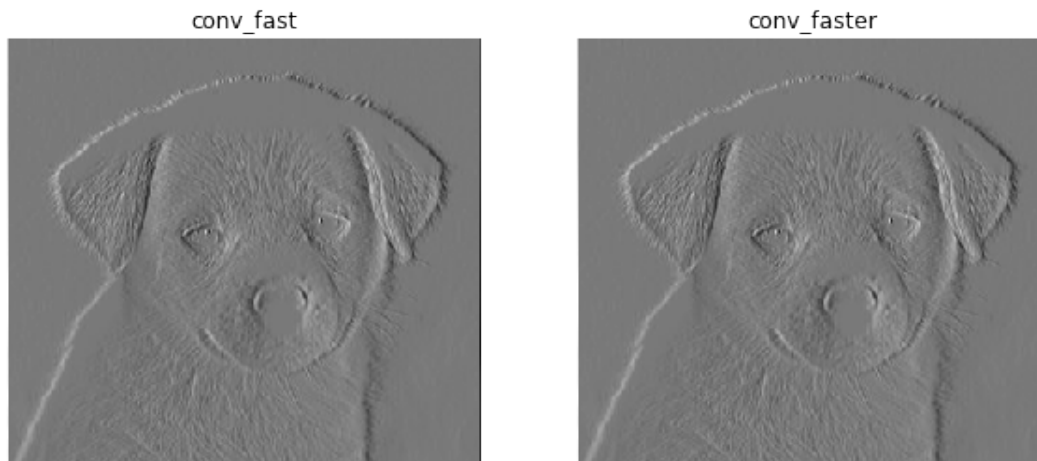
# Plot conv_nested output
plt.subplot(1,2,1)
plt.imshow(out_fast)
plt.title('conv_fast')
plt.axis('off')

# Plot conv_fast output
plt.subplot(1,2,2)
plt.imshow(out_faster)
plt.title('conv_faster')
plt.axis('off')

# Make sure that the two outputs are the same
if not (np.max(out_fast - out_faster) < 1e-10):
    print("Different outputs! Check your implementation.")

conv_fast: took 0.501648 seconds.
conv_faster: took 0.008152 seconds.

```



## 1.2 Part 2: Cross-correlation

Cross-correlation of two 2D signals  $f$  and  $g$  is defined as follows:

$$(f \star g)[m, n] = \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} f[i, j] \cdot g[i - m, j - n]$$

### 1.2.1 2.1 Template Matching with Cross-correlation (12 points)

Suppose that you are a clerk at a grocery store. One of your responsibilities is to check the shelves periodically and stock them up whenever there are sold-out items. You got tired of this laborious task and decided to build a computer vision system that keeps track of the items on the shelf.

Luckily, you have learned in CS131 that cross-correlation can be used for template matching: a template  $g$  is multiplied with regions of a larger image  $f$  to measure how similar each region is to the template.

The template of a product (`template.jpg`) and the image of shelf (`shelf.jpg`) is provided. We will use cross-correlation to find the product in the shelf.

Implement `cross_correlation` function in `filters.py` and run the code below.

- Hint: you may use the `conv_fast` function you implemented in the previous question.

```
In [11]: from filters import cross_correlation

# Load template and image in grayscale
img = io.imread('shelf.jpg')
img_grey = io.imread('shelf.jpg', as_grey=True)
temp = io.imread('template.jpg')
temp_grey = io.imread('template.jpg', as_grey=True)
```

```

# Perform cross-correlation between the image and the template
out = cross_correlation(img_grey, temp_grey)

# Find the location with maximum similarity
y,x = (np.unravel_index(out.argmax(), out.shape))

# Display product template
plt.figure(figsize=(25,20))
plt.subplot(3, 1, 1)
plt.imshow(temp)
plt.title('Template')
plt.axis('off')

# Display cross-correlation output
plt.subplot(3, 1, 2)
plt.imshow(out)
plt.title('Cross-correlation (white means more correlated)')
plt.axis('off')

# Display image
plt.subplot(3, 1, 3)
plt.imshow(img)
plt.title('Result (blue marker on the detected location)')
plt.axis('off')

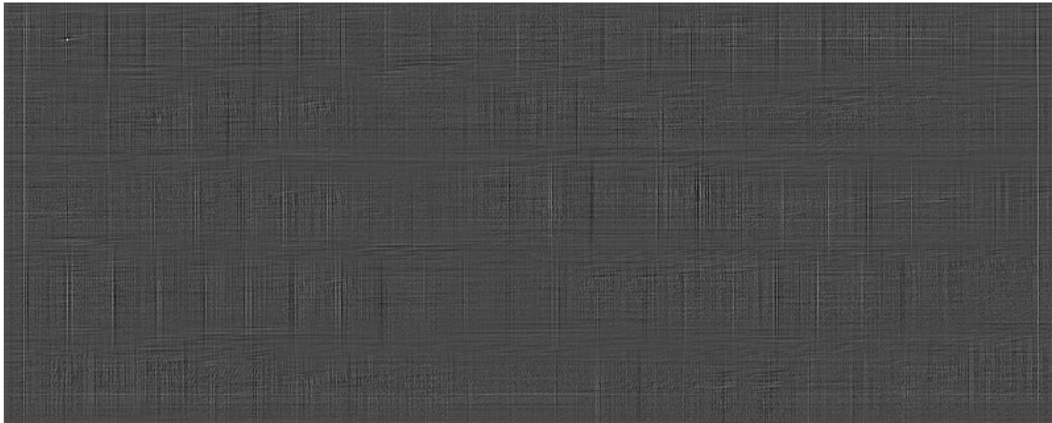
# Draw marker at detected location
plt.plot(x, y, 'bx', ms=40, mew=10)
plt.show()

```

Template



Cross-correlation (white means more correlated)



Result (blue marker on the detected location)



**Interpretation** How does the output of cross-correlation filter look like? Was it able to detect the product correctly? Explain what might be the problem with using raw template as a filter.

---

### 1.2.2 2.2 Zero-mean cross-correlation (6 points)

A solution to this problem is to subtract off the mean value of the template so that it has zero mean.

Implement `zero_mean_cross_correlation` function in `filters.py` and run the code below.

```
In [12]: from filters import zero_mean_cross_correlation

# Perform cross-correlation between the image and the template
out = zero_mean_cross_correlation(img_grey, temp_grey)

# Find the location with maximum similarity
y,x = (np.unravel_index(out.argmax(), out.shape))

# Display product template
plt.figure(figsize=(30,20))
plt.subplot(3, 1, 1)
plt.imshow(temp)
plt.title('Template')
plt.axis('off')

# Display cross-correlation output
plt.subplot(3, 1, 2)
plt.imshow(out)
plt.title('Cross-correlation (white means more correlated)')
plt.axis('off')

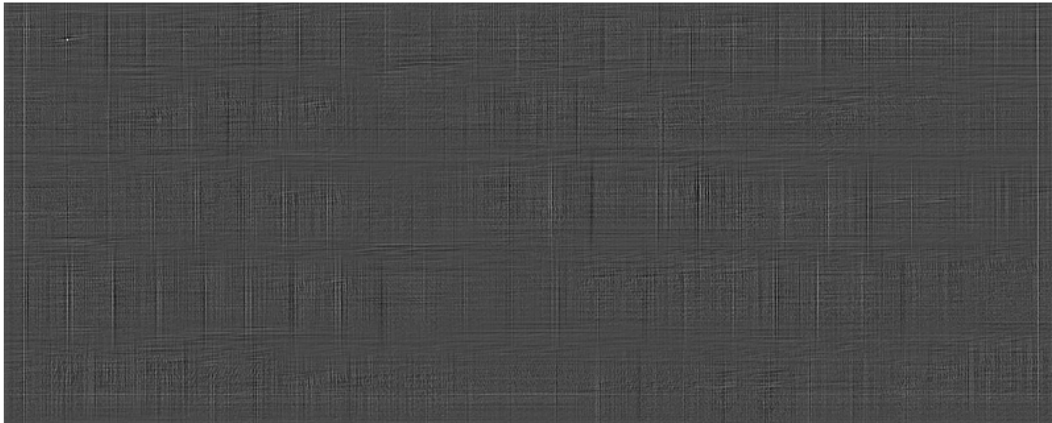
# Display image
plt.subplot(3, 1, 3)
plt.imshow(img)
plt.title('Result (blue marker on the detected location)')
plt.axis('off')

# Draw marker at detected location
plt.plot(x, y, 'bx', ms=40, mew=10)
plt.show()
```

Template



Cross-correlation (white means more correlated)



Result (blue marker on the detected location)



You can also determine whether the product is present with appropriate scaling and thresholding.

```
In [13]: def check_product_on_shelf(shelf, product):
```



```

out = zero_mean_cross_correlation(shelf, product)

# Scale output by the size of the template
out = out / float(product.shape[0]*product.shape[1])

# Threshold output (this is arbitrary, you would need to tune the threshold for a
out = out > 0.025

if np.sum(out) > 0:
    print('The product is on the shelf')
else:
    print('The product is not on the shelf')

# Load image of the shelf without the product
img2 = io.imread('shelf_soldout.jpg')
img2_grey = io.imread('shelf_soldout.jpg', as_grey=True)

plt.imshow(img)
plt.axis('off')
plt.show()
check_product_on_shelf(img_grey, temp_grey)

plt.imshow(img2)
plt.axis('off')
plt.show()
check_product_on_shelf(img2_grey, temp_grey)

```



The product is on the shelf



The product is not on the shelf

---

### 1.2.3 2.3 Normalized Cross-correlation (12 points)

One day the light near the shelf goes out and the product tracker starts to malfunction. The zero\_mean\_cross\_correlation is not robust to change in lighting condition. The code below demonstrates this.

```
In [29]: from filters import normalized_cross_correlation

# Load image
img = io.imread('shelf_dark.jpg')
img_grey = io.imread('shelf_dark.jpg', as_grey=True)
temp_grey = io.imread('template.jpg', as_grey=True)
# Perform cross-correlation between the image and the template
out = zero_mean_cross_correlation(img_grey, temp_grey)

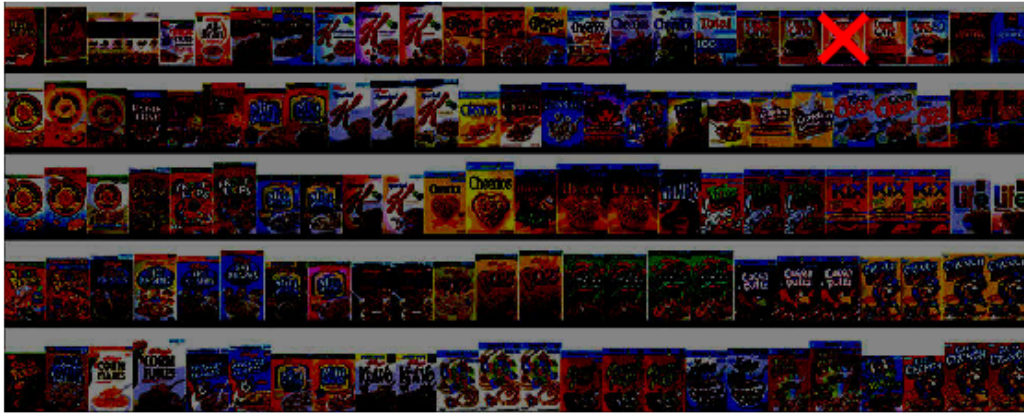
# Find the location with maximum similarity
y,x = (np.unravel_index(out.argmax(), out.shape))

# Display image
plt.imshow(img)
plt.title('Result (red marker on the detected location)')
plt.axis('off')

# Draw marker at detected location
plt.plot(x, y, 'rx', ms=25, mew=5)
plt.show()
```



Result (red marker on the detected location)



A solution is to normalize the pixels of the image and template at every step before comparing them. This is called **normalized cross-correlation**.

The mathematical definition for normalized cross-correlation of  $f$  and template  $g$  is:

$$(f \star g)[m, n] = \sum_{i, j} \frac{f[i, j] - \overline{f_{m, n}}}{\sigma_{f_{m, n}}} \cdot \frac{g[i - m, j - n] - \overline{g}}{\sigma_g}$$

where: -  $f_{m, n}$  is the patch image at position  $(m, n)$  -  $\overline{f_{m, n}}$  is the mean of the patch image  $f_{m, n}$  -  $\sigma_{f_{m, n}}$  is the standard deviation of the patch image  $f_{m, n}$  -  $\overline{g}$  is the mean of the template  $g$  -  $\sigma_g$  is the standard deviation of the template  $g$

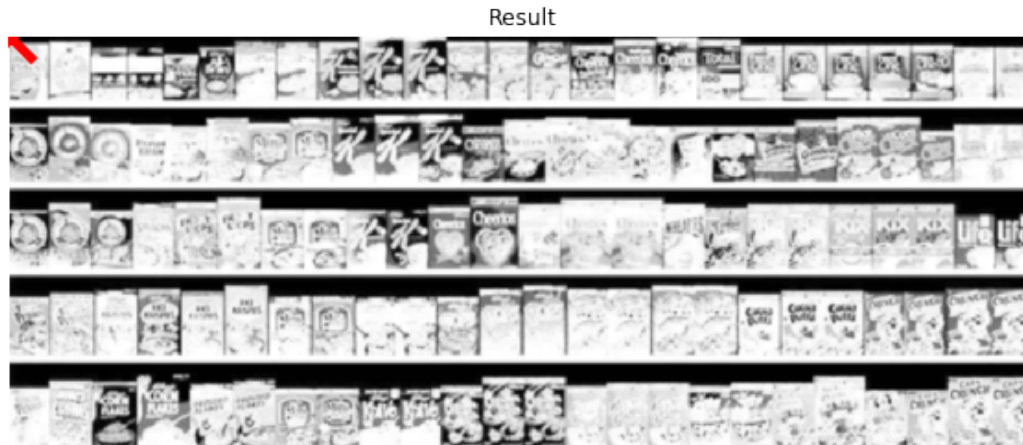
Implement `normalized_cross_correlation` function in `filters.py` and run the code below.

```
In [34]: from filters import normalized_cross_correlation
img_grey = io.imread('shelf_dark.jpg', as_grey=True)
temp_grey = io.imread('template.jpg', as_grey=True)

# Perform normalized cross-correlation between the image and the template
out = normalized_cross_correlation(img_grey, temp_grey)
# Find the location with maximum similarity
y,x=(np.unravel_index(out.argmax(), out.shape))
# Display image
plt.imshow(out)
plt.title('Result')
plt.axis('off')

# Draw marker at detected location
plt.plot(x, y, 'rx', ms=25, mew=5)
plt.show()
```

```
[[ 1.  1.  1.]
 [ 1.  2.  1.]
 [ 1.  1.  1.]]
[[-0.11111111 -0.11111111 -0.11111111]
 [-0.11111111 -0.11111111 -0.11111111]
 [-0.11111111 -0.11111111 -0.11111111]]
```



### 1.3 Part 3: Separable Filters

#### 1.3.1 3.1 Theory (10 points)

Consider a  $M_1 \times N_1$  image  $I$  and a  $M_2 \times N_2$  filter  $F$ . A filter  $F$  is **separable** if it can be written as a product of two 1D filters:  $F = F_1 F_2$ .

For example,

$$F = \begin{bmatrix} 1 & -1 \\ 1 & -1 \end{bmatrix}$$

can be written as a matrix product of

$$F_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix}, F_2 = \begin{bmatrix} 1 & -1 \end{bmatrix}$$

Therefore  $F$  is a separable filter.

Prove that for any separable filter  $F = F_1 F_2$ ,

$$I * F = (I * F_1) * F_2$$

**Your Answer:** Write your solution in this markdown cell. Please write your equations in [LaTeX equations](#).

Double click vào ô trống li.

### 1.3.2 3.2 Complexity comparison (10 points)

- (i) How many multiplications do you need to do a direct 2D convolution (i.e.  $I * F$ )?
- (ii) How many multiplications do you need to do 1D convolutions on rows and columns (i.e.  $(I * F_1) * F_2$ )?
- (iii) Use Big-O notation to argue which one is more efficient in general: direct 2D convolution or two successive 1D convolutions?

**Your Answer:** Write your solution in this markdown cell. Please write your equations in [LaTeX equations](#).

Double click vào ây tr li.

Now, we will empirically compare the running time of a separable 2D convolution and its equivalent two 1D convolutions. Gaussian kernel, widely used for blurring images, is one example of a separable filter. Run the code below to see its effect.

```
In [23]: # Load image
img = io.imread('dog.jpg', as_grey=True)

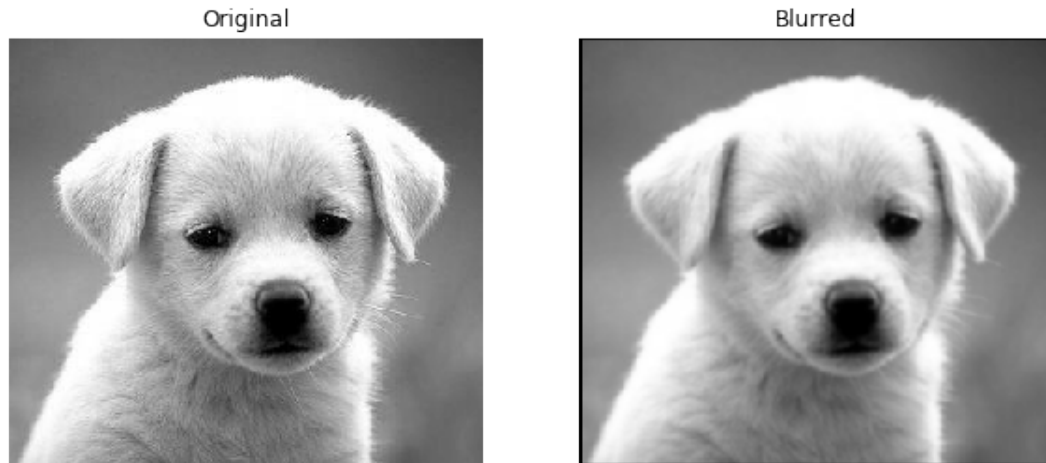
# 5x5 Gaussian blur
kernel = np.array(
[
    [1,4,6,4,1],
    [4,16,24,16,4],
    [6,24,36,24,6],
    [4,16,24,16,4],
    [1,4,6,4,1]
])

t0 = time()
out = conv_nested(img, kernel)
t1 = time()
t_normal = t1 - t0

# Plot original image
plt.subplot(1,2,1)
plt.imshow(img)
plt.title('Original')
plt.axis('off')

# Plot convolved image
plt.subplot(1,2,2)
plt.imshow(out)
plt.title('Blurred')
plt.axis('off')

plt.show()
```



In the below code cell, define the two 1D arrays (k1 and k2) whose product is equal to the Gaussian kernel.

```
In [25]: import math
         # The kernel can be written as outer product of two 1D filters
         k1 = None # shape (5, 1)
         k2 = None # shape (1, 5)

         ### YOUR CODE HERE
         u,s,v=np.linalg.svd(kernel)
         sqrt_value=math.sqrt(s[0])
         k1=u[:,0]*sqrt_value
         k2=v[:,0]*sqrt_value

         k1=np.array([k1]).T
         k2=np.array([k2])
         print(k1)
         print(k2)
         ### END YOUR CODE

         # Check if kernel is product of k1 and k2
         if not np.all(k1 * k2 == kernel):
             print('k1 * k2 is not equal to kernel')

         assert k1.shape == (5, 1), "k1 should have shape (5, 1)"
         assert k2.shape == (1, 5), "k2 should have shape (1, 5)"
```

```
[[ -1.]
 [ -4.]
 [ -6.]
 [ -4.]
```

```

[-1.]]
[[-1.          -0.68935454 -8.27797018  0.          -0.          ]]
k1 * k2 is not equal to kernel

```

We now apply the two versions of convolution to the same image, and compare their running time. Note that the outputs of the two convolutions must be the same.

```

In [271]: # Perform two convolutions using k1 and k2
          t0 = time()
          out_separable = conv_faster(img, k1)
          out_separable = conv_faster(out_separable, k2)
          t1 = time()
          t_separable = t1 - t0

          # Plot normal convolution image
          plt.subplot(1,2,1)
          plt.imshow(out)
          plt.title('Normal convolution')
          plt.axis('off')

          # Plot separable convolution image
          plt.subplot(1,2,2)
          plt.imshow(out_separable)
          plt.title('Separable convolution')
          plt.axis('off')

          plt.show()

          print("Normal convolution: took %f seconds." % (t_normal))
          print("Separable convolution: took %f seconds." % (t_separable))

```

Normal convolution



Separable convolution



Normal convolution: took 2.928843 seconds.  
Separable convolution: took 0.011029 seconds.

```
In [ ]: # Check if the two outputs are equal
        assert np.max(out_separable - out) < 1e-10
```